

Shravan M R

AI / LLM ENGINEER · GENERATIVE AI · AGENTIC SYSTEMS

✉ shravan.revanna@gmail.com +91 9945332995 🌐 [myselfshravan](https://myselfshravan.com) 📄 in/shravanrevanna 🌐 shravanrevanna.me

EDUCATION

Bachelor of Engineering - Computer Science and Engineering (AI & ML)

2021 – 2025

M S Ramaiah Institute of Technology, Bangalore

Current CGPA: 8.81/10

Coursework: Data Structures and Algorithms, Design and Analysis of Algorithms, Database Management Concepts and Systems, Operating Systems, Deep Learning, Machine Learning, Object Oriented Programming using Java

TECHNICAL SKILLS

AI / ML: LLMs (GPT, Llama, Mistral), Generative AI, RAG, Fine-tuning, Prompt Engineering, Agentic / Multi-agent Systems, MCP, text2sql, NLP, PyTorch, CLIP

Languages: Python, Java, Kotlin, TypeScript

Frameworks and Libraries: PyTorch, FastAPI, LiteLLM, ChromaDB, Node.js, React.js, Next.js, Ktor

Data Management: Relational (PostgreSQL), Non-Relational (Firestore, Redis), Vector DBs (ChromaDB)

Cloud & MLOps: AWS (Lambda, EC2), Azure, Docker, Git, Postman, Mixpanel, Appsflyer

EXPERIENCE

Klydo | Product Engineer

Jan 2026 – Present

- Built **Sherlock**, an **LLM-based fraud-detection AI agent** for Take-Now-Pay-Later, flagging high-risk orders at checkout with **96% precision**; enabled a safer pay-later rollout and prevented revenue leakage without throttling good buyers.
- Designed and deployed an **agentic data layer** (Lexi + DB Explorer MCP) using **text2sql** + **LLM tool-calling**, letting non-technical founders query production **PostgreSQL** in plain English over WhatsApp/Slack; eliminated dependency on dashboards.
- Engineered an **atomic, concurrency-safe referral & wallet ledger**; scaled referral volume **4.7x** (80 to 378/week) and processed **400+ paid orders** with audit-safe transactions.
- Shipped **coupon hub & TNPL pricing logic** (+5pp checkout conversion, 82% → 87%) and **growth infra**: automated ingestion pipelines and unified **S2S tracking** across GA4, Mixpanel, Appsflyer, and Clevertap.

Graphene AI | Machine Learning Engineer

Sep 2025 – Dec 2025

- Designed an **agentic AI pipeline for LLM evaluation and deployment** within a distributed microservices architecture.
- Built an in-house **deploy console** and maintained a **LiteLLM reverse-proxy** gateway routing LLM traffic across teams.
- Architected and restructured **microservices** for high-availability workflows, embedding **gen-AI** into core business processes.

Udaan (incl. Udaan Capital) | SWE Intern

Jan 2025 – Aug 2025

- Enhanced **Udaan Capital's ML risk-evaluation engine**, improving onboarding accuracy and reducing **feature drift** across model retrains.
- Built full-stack features across internal teams using **Kotlin** (backend) and **TypeScript/React** (frontend).
- Automated **competitor-data ingestion** pipelines, cutting setup errors **70%** and doubling onboarding throughput, with dynamic request signing and incremental sync over volatile external sources.

CADRE ODR | ML Intern

Jan 2024 – Mar 2024

- Developed a real-time **grievance redressal** (RAG) application using ChromaDB and **Mistral 7B & Llama 2** LLMs for user interaction and information extraction.
- Fine-tuned a **BART model** on a custom grievance dataset for zero-shot **hierarchical classification** of government departments. Integrated the classification model with chatbot flow to guide users through a hierarchical grievance path.

PROJECTS

ParseFi - AI Bank Statement Analyzer | React, Supabase/PostgreSQL, LLM tool-calling, text2sql, MCP

2025

- Parses and consolidates bank statements into a financial-intelligence layer with running balances, reconciliation, and an **MCP-connected AI chatbot** answering spending questions in plain English via **LLM tool-calling** + **text2sql**.
- Engineered $O(\log n)$ B-tree date-range queries, hash-based dedup, and sliding-window batch processing for auto-reconciliation at scale.

sudo-trade - Autonomous Multi-Agent Trading System | Multi-agent orchestration, LLMs, real-time data

2026

- Built **6 autonomous AI agents** that research, debate, and trade Indian stocks; a bull and bear agent argue **adversarially** before each trade decision.
- Ran **483 debates** and **148+ trades** across **370K tokens** for under a dollar a day - agentic design that stays cost-disciplined.

Enigma Escape - LLM Prompt-Engineering Game | Python, Streamlit, Llama-3 70B, Firebase

Dec 2023

- A linguistic puzzle game where players coax an **LLM** into saying exact target phrases via indirect prompting - effectively **red-teaming** model failure modes.
- Integrated **Firebase** auth and progression tracking to capture activity and score players across escalating levels.